

# Solving Lunar Lander using Deep Reinforcement Learning

AI course

Lemuel Puglisi, UniCT - 2023




Università  
di Catania

# The problem

Design a **learning agent** that learns how to land the rocket inside the landing pad, by turning engines on and off.

Environment provided by:



# The environment

- Partially observable
- Deterministic
- Sequential
- Static
- Continuous
- Single-agent

# Action space

The action space is discrete:

- 0: do nothing
- 1: fire left orientation engine
- 2: fire main engine
- 3: fire right orientation engine

# Observation space

The state is an 8-dimensional vector: the coordinates of the lander in  $(x, y)$ , its linear velocities in  $(\dot{x}, \dot{y})$ , its angle, its angular velocity, and two booleans that represent whether each leg is in contact with the ground or not.

# Starting state

The lander starts at the top center of the viewport with a random initial force applied to its center of mass.

# Rewards mechanism (1)

- decreased proportionally to the distance to the landing pad
- decreased proportionally to the speed of the lander
- decreases proportionally to the angle w.r.t. the ground
- increased by 10 points for each leg in contact with the ground
- decreased by 0.03 when side engines are actioned
- decreased by 0.3 when the main engine is actioned

# Reward mechanism (2)

The episode ends if the lander crashes or gets outside of the viewport.  
When an episode ends, the agent:


- receives an additional reward of +100 for landing safely
- receives an additional penalty (negative reward) of -100 for crashing the lander.



# How does the agent learn?

# Reinforcement learning

At time  $t$ , the agent interacts with the environment, which has a state  $s_t$ , by performing an action  $a_t$ . The agent receives a reward  $r_t$  based on the pair  $(s_t, a_t)$  and the environment changes to a new state  $s_{t+1}$ .



# Markov decision process

This process can be formalized as a Markov Decision Process  $(S, A, P, R, \gamma)$  where:

- $S$  is the set of the environment states
- $A$  is the set of possible actions
- $R$  is the reward distribution
- $P$  is the transition distribution

# Policy

The action  $a_t$  performed by the agent is determined by a function  $\pi : S \rightarrow A$  called **policy**.

# Return

We want to find an optimal policy  $\pi^*$ , i.e. a policy that optimizes the return  $R_{t_0}$

$$R_t = \sum_{t'=t}^T \gamma^{(t'-t)} r_{t'}$$

Where  $\gamma \in [0, 1]$  is called **discount rate** and is used to balance the trade-off between short-term and long-term rewards, and  $T$  is the total number of steps.

# Optimal Q-value function

Let  $Q^* : S \times A \rightarrow \mathbb{R}$  be a function, called optimal Q-value function, that predicts the final return we will receive by choosing an action  $a_t$  given a state  $s_t$  and then proceeding with an optimal policy  $\pi^*$ . The  $Q^*$  function satisfies the **Bellman equation**, thus can be written as:

$$Q^*(s, a) = r + \gamma \max_{a'} Q^*(s', a')$$

Where  $s'$  is the next state given  $(s, a)$ .

# Optimal policy

Given the Q-value function  $Q^*$ , defining the optimal policy is trivial:

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

# Deep Q-learning

Deep Q-learning tries to estimate the  $Q^*(s, a)$  function using a parametrized function  $Q(s, a; \theta)$  (i.e., a deep neural network). If we obtain  $Q(s, a; \theta) \approx Q^*(s, a)$ , then the agent policy will be set to:

$$\pi(s) = \arg \max_a Q(s, a; \theta)$$

The algorithm can be divided in two phases, **collecting experience** and **learning**, repeated until convergence.



# Collecting experience

Collect experience  $(s, a, r, s')$  in a memory, called **replay buffer**  $D$ , by playing the agent for  $m$  episodes in the simulator.

# $\epsilon$ -greedy policy

To perform better exploration, we adopt an  $\epsilon$ -greedy policy:

$$\pi_{\epsilon}(s) = \begin{cases} \text{random action } a & \text{with probability } \epsilon \\ \arg \max_{a \in A} Q(s, a; \theta_i) & \text{with probability } 1 - \epsilon \end{cases}$$

The  $\epsilon$  value can be scheduled during the epochs, starting from a high value and decreasing it to strengthen the agent's strategy.

# Learning phase (1)

Sample a batch  $B \subset D$  of random experience  $(s, a, r, s') \in B$  from the replay buffer  $D$  and computing the **temporal difference error**  $\mathcal{L}$ :

$$\mathcal{L}(s, a, r, s', \theta_t) = Q(s, a; \theta_t) - (r + \gamma \max_a Q(s', a; \theta_{t-1}))$$

By minimizing  $\mathcal{L}$  we **force** the Bellman equation:

$$0 = Q(s, a, \theta_t) - (r + \gamma \max_a Q(s', a; \theta_{t-1}))$$

$$Q(s, a, \theta_t) = (r + \gamma \max_a Q(s', a; \theta_{t-1}))$$

# Deep Q-learning (3)

The parameters are optimized by minimizing the following loss function computed on the batch:

$$\mathcal{J}(\theta_t) = \frac{1}{|B|} \sum_{(s,a,r,s') \in B} \mathcal{L}(s, a, r, s', \theta_t)$$

E.g. by using common optimizers like SGD.

# Avoiding catastrophic forgetting (1)

Did you notice that we use two different versions of  $Q$  to compute the loss?

$$\mathcal{L}(s, a, r, s', \theta_t) = Q(s, a; \theta_t) - (r + \gamma \max_a Q(s', a; \theta_{t-1}))$$

$Q(\cdot, \cdot; \theta)$  is called the Q-network, while  $Q(\cdot, \cdot; \theta_{t-1})$  is called the target network, and it is updated every  $\tau$  batches with the weights of the Q-network.

# Avoiding catastrophic forgetting (2)

Using two networks mitigates the catastrophic forgetting phenomenon, i.e., forgetting skills learned in the past while learning new skills.

**Let's get our hands dirty**



Università  
di Catania

# Implementing the agent

The following pseudo-code gives an idea of how the agent is implemented:

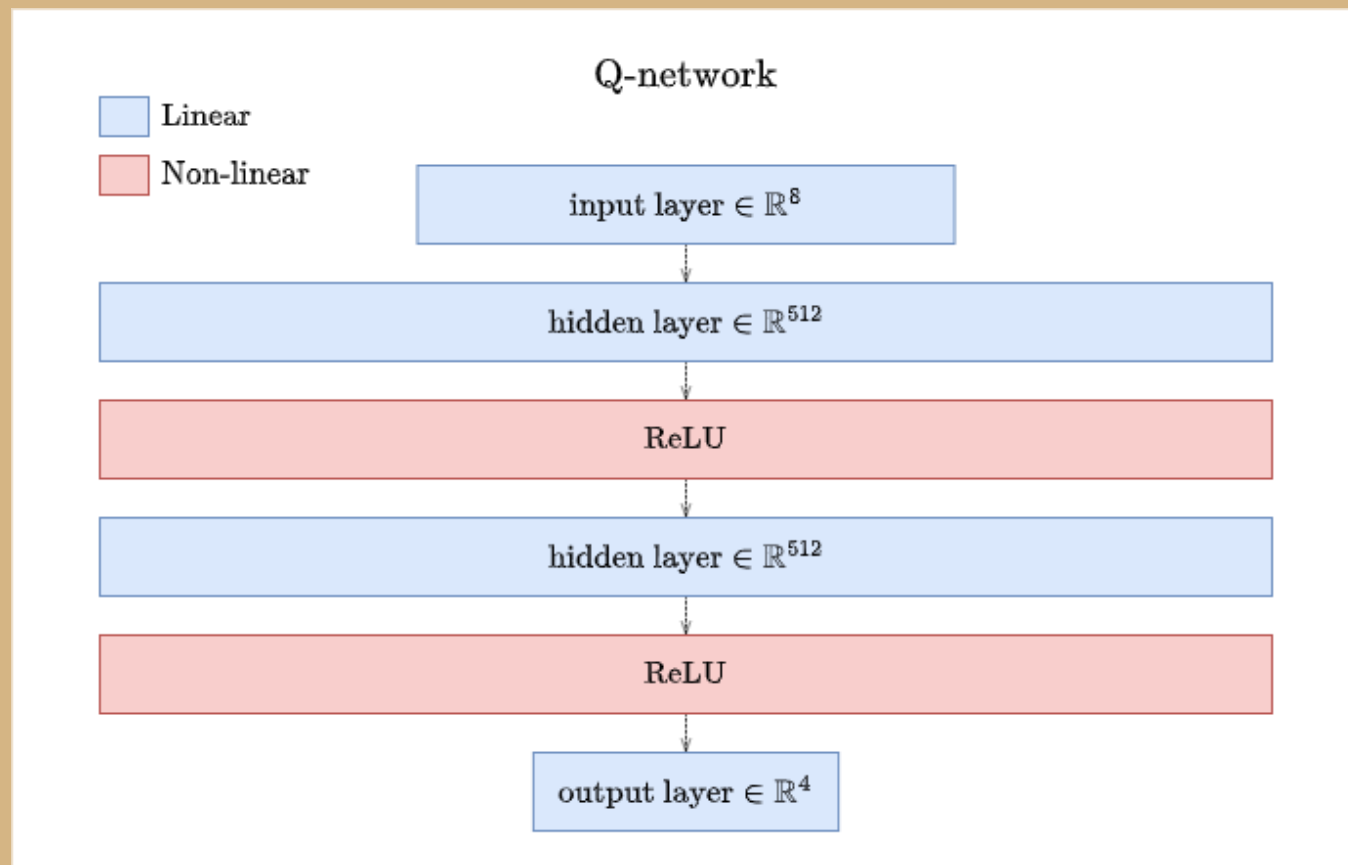
```
class LearningAgent:

    def __init__(self, q_network):
        self.q_network = q_network

    def policy(self, observations):
        best_action = argmax(self.q_network(observations))
        return best_action
```



# Q-network architecture



# $\epsilon$ -scheduling

The  $\epsilon$  value of the  $\epsilon$ -greedy policy is scheduler w.r.t. the current epoch  $t$  using the following function:

$$\epsilon(t) = \max \left( 0.1, 1 - \frac{2t}{E} \right)$$

Where  $E$  is the total number of epochs.

# Training settings

- Optimizer: Adam
- Learning rate: 0.001
- Replay memory size:  $5 \times 10^5$
- Sync-rate  $\tau$ : every 10 batches
- Batch size: 512


# Winning the game

The game is considered to be won when the agent achieves a return of around 200. Our agent **surpassed this threshold** after approximately 35 epochs.

# Ablation study

# Ablation study: Capacity


The following figure shows the performance of the trained agent using a replay memory capacity of  $5 \times 10^5$ ,  $5 \times 10^4$ , and  $5 \times 10^3$ . We note that (1) performance increases proportionally with capacity; (2) if the capacity is not sufficient (red line), the agent's performance may degrade over time instead of improving (catastrophic forgetting).



# Ablation study: Sync-rate $\tau$

The training is repeated using different target network update frequencies. Generally, we observe that with a higher frequency, such as 1 update per batch (red line), there is a slower progression. By decreasing the update frequency, for example, 1 every 10 or 100 batches (blue or green line), we notice a faster progression and greater stability.





# Conclusions

The proposed algorithm can solve the Lunar Lander problem by learning an optimal strategy in a limited number of iterations. However, it also presents some limitations, such as difficulty in defining the reward function and managing the algorithm's hyperparameters, which play a fundamental role. Currently, Deep Reinforcement Learning is becoming more widely used due to its flexibility in addressing a variety of problems.

# Open-Source Implementation

[Code](#) [Issues](#) [Pull requests](#) [Actions](#) [Projects](#) [Security](#) [Insights](#)

main 1 branch 0 tags

Go to file

Add file

Code



lemuelatQSA doc: add play example

39ef354 last week 12 commits



docs

doc: update GIF

2 weeks ago



models

feat: add models checkpoints

2 weeks ago



src/aicourse

doc: update GIF

2 weeks ago



.gitignore

update gitignore

2 weeks ago



LICENSE.md

init project

2 weeks ago



README.md

doc: add play example

last week



play.py

fix: run play also on cpu

last week



pyproject.toml

fix: add box2D, typo in torch

last week



run\_experiments.sh

init project

2 weeks ago



train.py

init project

2 weeks ago



README.md

# Lunar Lander - DQN

## About

Simple PyTorch implementation of Deep Q-learning Algorithm to play Lunar Lander.

[reinforcement-learning](#) [deep-q-learning](#)  
[lunar-lander](#) [open-ai-gym](#)

Readme

MIT license

5 stars

1 watching

0 forks

## Releases

No releases published

## Packages

No packages published

## Languages

Python 94.4% Shell 5.6%

# Training the agent

```
(your_env) <Lunar-Lander-DQN> python train.py --help
```

```
usage: train.py [-h] [--dest DEST] [--epochs EPOCHS] [--episodes EPISODES] [--batch-size BATCH_SIZE] [--capacity CAPACITY]
               [--sync-rate SYNC_RATE]
```

optional arguments:

-h, --help	show this help message and exit
--dest DEST	destination folder
--epochs EPOCHS	number of epochs
--episodes EPISODES	number of episodes to play <b>in</b> an epoch
--batch-size BATCH_SIZE	batch size on trainin phase
--capacity CAPACITY	capacity of the replay memory
--sync-rate SYNC_RATE	sync rate of the target network

# Agent in action

```
(your_env) <Lunar-Lander-DQN> python play.py --help
```

```
usage: play.py [-h] [--model-ckpt MODEL_CKPT] [--episodes EPISODES]
```

```
optional arguments:
```

```
-h, --help            show this help message and exit  
--model-ckpt MODEL_CKPT  
--episodes EPISODES
```



**Thank you!**

